

## DATA ENGINEERING TOOLS – SEGREGATED

Segment	General	AWS
<b>Data Ingestion</b>	<b>Ingestion Tools:</b> Apache Kafka Apache NiFi AWS Kinesis Logstash	<b>AWS Glue:</b> Use Glue Crawlers to discover and catalog metadata from various data sources. Glue ETL jobs for transforming and loading data.  <b>Amazon Kinesis:</b> Kinesis Data Streams for real-time data streaming. Kinesis Data Firehose for loading streaming data into data stores.  <b>AWS DataSync:</b> Transfer data from on-premises to AWS.
<b>Data Storage</b>	<b>Data Warehouses:</b> Amazon Redshift Google BigQuery Snowflake  <b>Data Lakes:</b> Amazon S3 Azure Data Lake Storage Google Cloud Storage  <b>Databases:</b> PostgreSQL MySQL MongoDB Cassandra	<b>Amazon S3:</b> As a data lake for storing raw and processed data. Versioning and lifecycle policies for managing data.  <b>Amazon Redshift:</b> For data warehousing and complex queries.  <b>Amazon DynamoDB:</b> For NoSQL database requirements.
<b>Data Processing</b>	<b>Batch Processing:</b> Apache Spark Apache Flink Hadoop MapReduce  <b>Stream Processing:</b> Apache Kafka Streams Apache Storm Apache Flink  <b>ETL (Extract, Transform, Load):</b> Apache Beam Apache Airflow Talend	<b>Amazon EMR (Elastic MapReduce):</b> For big data processing using frameworks like Apache Spark and Hadoop.  <b>AWS Glue:</b> Serverless ETL service for data transformation and preparation.  <b>AWS Lambda:</b> For serverless event-driven processing.
<b>Data Transformation</b>	<b>Data Preparation:</b> Pandas (Python library) Apache Beam  <b>Data Cleansing:</b> Trifacta OpenRefine  <b>Data Masking/Anonymization:</b> Google DLP Apache Nifi	<b>AWS Glue:</b> Use Glue jobs for ETL transformations.  <b>AWS Step Functions:</b> Orchestrate and coordinate multiple AWS services in a serverless workflow.

<b>Analytics and Reporting</b>	<p><b>Business Intelligence Tools:</b> Tableau Power BI Looker</p> <p><b>Analytics Platforms:</b> Databricks Google Analytics Mixpanel</p>	<p><b>Amazon QuickSight:</b> Business intelligence service for visualizing and analyzing data.</p> <p><b>Amazon Athena:</b> Serverless query service for analyzing data in Amazon S3.</p>
<b>Data Orchestration</b>	<p><b>Workflow Management:</b> Apache Airflow Luigi Prefect</p> <p><b>Job Scheduling:</b> Cron Apache Oozie</p>	<p><b>Apache Airflow on Amazon MWAA (Managed Workflows for Apache Airflow):</b> Orchestrate and schedule complex data workflows.</p> <p><b>AWS Step Functions:</b> For serverless workflow orchestration.</p>
<b>Monitoring and Logging</b>	<p><b>Logging:</b> ELK Stack (Elasticsearch, Logstash, Kibana) Splunk</p> <p><b>Monitoring:</b> Prometheus Grafana</p>	<p><b>Amazon CloudWatch:</b> For monitoring AWS resources and applications.</p> <p><b>AWS CloudTrail:</b> For logging AWS API calls.</p>
<b>Data Data Quality and Governance</b>	<p><b>Data Quality Tools:</b> Informatica Talend Apache Griffin</p> <p><b>Metadata Management:</b> Collibra Apache Atlas</p>	<p><b>AWS Glue DataBrew:</b> For data profiling, cleaning, and exploration.</p> <p><b>AWS Lake Formation:</b> Set up and enforce security, governance, and auditing policies.</p>
<b>Security and Access Control</b>	<p><b>Encryption:</b> TLS/SSL HDFS Encryption</p> <p><b>Access Control:</b> Apache Ranger AWS IAM Google Cloud Identity and Access Management (IAM)</p>	<p><b>AWS IAM (Identity and Access Management):</b> Manage access to AWS resources.</p> <p><b>AWS Key Management Service (KMS):</b> Encrypt data at rest and in transit.</p>
<b>Data Science Integration</b>	<p><b>Model Deployment:</b> TensorFlow Serving MLflow PMML (Predictive Model Markup Language)</p> <p><b>Notebook Environments:</b> Jupyter Notebooks Google Colab Databricks Notebooks</p>	<p><b>Amazon SageMaker:</b> For building, training, and deploying machine learning models.</p>
<b>Architectural Patterns</b>	<p><b>Lambda Architecture:</b> Combines batch and stream processing for real-time and batch processing.</p>	<p><b>Serverless Architecture:</b> Leverage services like Lambda, Glue, and Step Functions for serverless processing.</p>

	<b>Kappa Architecture:</b> Simplifies the Lambda Architecture using only stream processing.	<b>Data Lake Architecture:</b> Utilize S3 as a central data lake to store structured and unstructured data.
<b>Data Versioning and Lineage</b>	<b>Version Control:</b> Git DVC (Data Version Control)  <b>Lineage Tracking:</b> Apache Atlas DataHub	
<b>Cloud Integration</b>	<b>Cloud Platforms:</b> AWS, Azure, Google Cloud Platform (GCP)  <b>Serverless Computing:</b> AWS Lambda Azure Functions Google Cloud Functions	<b>AWS Direct Connect or VPN:</b> Connect on-premises data centers to AWS.  <b>AWS SDKs and CLI:</b> Integrate and automate AWS services using SDKs and the Command Line Interface.

Segment	Microsoft Azure	Google Cloud Platform
<b>Data Ingestion</b>	<b>Azure Data Factory:</b> Orchestrate and automate data workflows. Support for data movement from various sources to data lakes or warehouses.  <b>Azure Event Hubs:</b> Ingest and process massive amounts of streaming data.	<b>Cloud Pub/Sub:</b> Real-time messaging service for event-driven architectures.  <b>Cloud Storage:</b> Object storage for batch uploads.
<b>Data Storage</b>	<b>Azure Data Lake Storage:</b> Scalable and secure data lake storage.  <b>Azure SQL Data Warehouse (now part of Azure Synapse Analytics):</b> Enterprise-grade analytics service.  <b>Azure Cosmos DB:</b> Globally distributed, multi-model database for operational and analytical workloads.	<b>BigQuery:</b> Fully-managed, serverless data warehouse for analytics.  <b>Cloud Storage:</b> Object storage for raw data and backups.  <b>CloudSQL:</b> Managed relational databases.
<b>Data Processing</b>	<b>Azure Databricks:</b> Apache Spark-based analytics platform for big data and machine learning.  <b>HDInsight:</b> Fully managed cloud service for big data analytics using Hadoop, Spark, HBase, and more.  <b>Azure Stream Analytics:</b> Real-time analytics on streaming data.	<b>Dataflow:</b> Fully managed stream and batch processing using Apache Beam.  <b>Dataprep by Trifacta:</b> Cloud-native data preparation service.  <b>Dataproc:</b> Managed Apache Spark and Hadoop service.
<b>Data Transformation</b>	<b>Azure Data Factory:</b> Transform and clean data using data flows and transformations.  <b>Azure HDInsight:</b> Leverage Apache Spark or Hive for data transformation.	<b>Dataflow:</b> Apache Beam for ETL pipelines.  <b>Cloud Dataprep:</b> Visual data preparation tool.

<b>Analytics and Reporting</b>	<p><b>Power BI:</b> Business Intelligence and visualization.</p> <p><b>Azure Synapse Studio:</b> Integrated analytics and data exploration.</p>	<p><b>BigQuery:</b> For ad-hoc queries and analytics.</p> <p><b>Looker, Tableau, or Data Studio:</b> Business intelligence and visualization tools.</p>
<b>Data Orchestration</b>	<p><b>Azure Data Factory:</b> Schedule and orchestrate data workflows.</p> <p><b>Azure Logic Apps:</b> Automate workflows and integrate services, including data services.</p>	<p><b>Cloud Composer:</b> Managed Apache Airflow for workflow orchestration.</p> <p><b>Cloud Scheduler:</b> Fully managed cron job scheduler.</p>
<b>Monitoring and Logging</b>	<p><b>Azure Monitor:</b> Monitor the performance and health of resources.</p> <p><b>Azure Log Analytics:</b> Collect and analyze log data.</p>	<p><b>Cloud Monitoring:</b> Infrastructure and application monitoring.</p> <p><b>Cloud Logging:</b> Centralized log management.</p>
<b>Data Data Quality and Governance</b>	<p><b>Azure Purview:</b> Unified data governance service for discovering, understanding, and managing data.</p> <p><b>Azure Data Catalog:</b> Discover, register, and manage data asset.</p>	<p><b>Cloud Data Catalog:</b> Fully managed and scalable metadata management service.</p> <p><b>Cloud Data Loss Prevention (DLP):</b> Sensitive data discovery and redaction.</p>
<b>Security and Access Control</b>	<p><b>Azure Active Directory (AAD):</b> Identity and access management.</p> <p><b>Azure Key Vault:</b> Securely store and manage sensitive information like keys and secrets.</p>	<p><b>Cloud Identity and Access Management (IAM):</b> Access control for GCP resources.</p> <p><b>Cloud Key Management Service (KMS):</b> Manage cryptographic keys.</p>
<b>Data Science Integration</b>	<p><b>Azure Machine Learning:</b> End-to-end platform for building, training, and deploying machine learning models.</p>	<p><b>AI Platform:</b> Managed services for building, training, and deploying machine learning models.</p> <p><b>Notebooks:</b> AI Platform Notebooks or Jupyter Notebooks on AI Platform.</p>
<b>Architectural Patterns</b>	<p><b>Modern Data Warehouse (Azure Synapse Analytics):</b> Combines big data and data warehousing for analytics.</p> <p><b>Event-Driven Architectures:</b> Use Azure Event Hubs and Azure Functions for event-driven processing.</p>	<p><b>Serverless Architecture:</b> Utilize serverless services like Cloud Functions.</p> <p><b>Data Lake and Data Warehouse:</b> Combine Cloud Storage and BigQuery for cost-effective storage and analytics.</p>
<b>Data Versioning and Lineage</b>		<p><b>Cloud Data Catalog:</b> Track and manage data lineage.</p> <p><b>BigQuery:</b> Keep track of changes with versioned tables.</p>
<b>Cloud Integration</b>	<p><b>Azure Functions:</b> Serverless computing for event-driven solutions.</p> <p><b>Azure Logic Apps:</b> Connect and automate workflows across cloud and on-premises services.</p>	<p><b>Cloud Functions:</b> Serverless computing for event-driven functions.</p> <p><b>Cloud Run:</b> Fully managed compute platform for containerized applications.</p>